

# Proximal Neural Networks: Marrying Variational Methods and Artificial Intelligence

## III – Neural networks

Audrey REPETTI<sup>†</sup>, Nelly PUSTELNIK<sup>◊</sup>, Jean-Christophe PESQUET\*

\* CentraleSupélec, Université Paris-Saclay, Gif-sur-Yvettes, France

◊ CNRS, ENS Lyon, Lyon, France

<sup>†</sup> Heriot-Watt University & Maxwell Institute for Mathematical Sciences, Edinburgh, UK

TUTORIAL – EUSIPCO 2025 – Palermo, Italy

# Unified framework

## Inference framework: feed-forward NN

$$\begin{aligned} (\forall \mathbf{x}^{[0]} \in \mathbb{R}^{N_0}) \quad \mathbf{x}^{[K]} &= \mathfrak{L}_{\Theta}^K(\mathbf{x}^{[0]}) \\ &= \mathfrak{T}_{\Theta_K} \circ \dots \circ \mathfrak{T}_{\Theta_1}(\mathbf{x}^{[0]}), \end{aligned}$$

### Layer/iteration

$$\mathfrak{T}_{\Theta_k}: \mathbb{R}^{N_{k-1}} \rightarrow \mathbb{R}^{N_k}: \mathbf{x} \mapsto \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + \mathbf{b}_k),$$

- ▶  $\mathbf{L}_k: \mathbb{R}^{N_{k-1}} \rightarrow \mathbb{R}^{N_k}$ : linear operator,
- ▶  $\mathbf{b}_k \in \mathbb{R}^{N_k}$ : shift parameter,
- ▶  $\mathfrak{D}_{\Lambda_k}: \mathbb{R}^{N_k} \rightarrow \mathbb{R}^{N_k}$ : nonlinear operator parametrized by  $\Lambda_k$ .

**Parameters:**  $\Theta = \cup_{k=1}^K \Theta_k$  with  $\Theta_k = \{\Lambda_k, \mathbf{L}_k, \mathbf{b}_k\}$ .

# Training

DATASET :  $\mathcal{S} = \{(\bar{\mathbf{x}}_j, \mathbf{z}_j) \in \mathbb{R}^N \times \mathbb{R}^M \mid j \in \mathbb{I} \cup \mathbb{J}\}$  pairs of (groundtruth, observation)

- *training set*  $(\bar{\mathbf{x}}_j, \mathbf{z}_j)_{j \in \mathbb{I}}$  with size  $|\mathbb{I}|$
- *testing set*  $(\bar{\mathbf{x}}_j, \mathbf{z}_j)_{j \in \mathbb{J}}$  with size  $|\mathbb{J}|$

EXAMPLE:  $\mathbf{z}_j = \mathbf{A}\bar{\mathbf{x}}_j + \varepsilon_j$



$\bar{\mathbf{x}}_\ell$



$\mathbf{z}_\ell$

# Training

DATASET :  $\mathcal{S} = \{(\bar{\mathbf{x}}_j, \mathbf{z}_j) \in \mathbb{R}^N \times \mathbb{R}^M \mid j \in \mathbb{I} \cup \mathbb{J}\}$  pairs of (groundtruth, observation)

- *training set*  $(\bar{\mathbf{x}}_j, \mathbf{z}_j)_{j \in \mathbb{I}}$  with size  $|\mathbb{I}|$
- *testing set*  $(\bar{\mathbf{x}}_j, \mathbf{z}_j)_{j \in \mathbb{J}}$  with size  $|\mathbb{J}|$

TRAINING STAGE: find  $\hat{\Theta} \in \operatorname{Argmin}_{\Theta} \frac{1}{|\mathbb{I}|} \sum_{j \in \mathbb{I}} \ell(\bar{\mathbf{x}}_j, \mathfrak{L}_{\Theta}^K(\mathbf{z}_j)).$

- $\ell$ : loss function
- High-dimensional non-convex optimization problem
- Use of auto-differentiation in standard toolboxes (e.g., Pytorch).

# Training

DATASET :  $\mathcal{S} = \{(\bar{\mathbf{x}}_j, \mathbf{z}_j) \in \mathbb{R}^N \times \mathbb{R}^M \mid j \in \mathbb{I} \cup \mathbb{J}\}$  pairs of (groundtruth, observation)

- *training set*  $(\bar{\mathbf{x}}_j, \mathbf{z}_j)_{j \in \mathbb{I}}$  with size  $|\mathbb{I}|$
- *testing set*  $(\bar{\mathbf{x}}_j, \mathbf{z}_j)_{j \in \mathbb{J}}$  with size  $|\mathbb{J}|$

TRAINING STAGE: find  $\hat{\Theta} \in \operatorname{Argmin}_{\Theta} \frac{1}{|\mathbb{I}|} \sum_{j \in \mathbb{I}} \ell(\bar{\mathbf{x}}_j, \mathfrak{L}_{\Theta}^K(\mathbf{z}_j)).$

- $\ell$ : loss function
- High-dimensional non-convex optimization problem
- Use of auto-differentiation in standard toolboxes (e.g., Pytorch).

EVALUATION: A properly trained network must satisfy

$$(\forall j \in \mathbb{J}) \quad \bar{\mathbf{x}}_j \approx \mathfrak{L}_{\Theta}^K(\mathbf{z}_j)$$

# NN example: Multi-Layer-Perceptron (MLP)

**DEFINITION:** [Scaman, Virmaux, 2018]

A  $K$ -layer MLP is of the form of

$$\mathfrak{L}_{\Theta}^K = \mathfrak{T}_{\Theta_1} \circ \dots \circ \mathfrak{T}_{\Theta_K} \quad \text{with} \quad \mathfrak{T}_{\Theta_k} = \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + \mathbf{b}_k),$$

where  $(\mathbf{L}_k)_{1 \leq k \leq K}$  are dense matrices (i.e., fully connected layers).

**REMARK:** When  $K = 1$ , the network is then referred to as a perceptron.

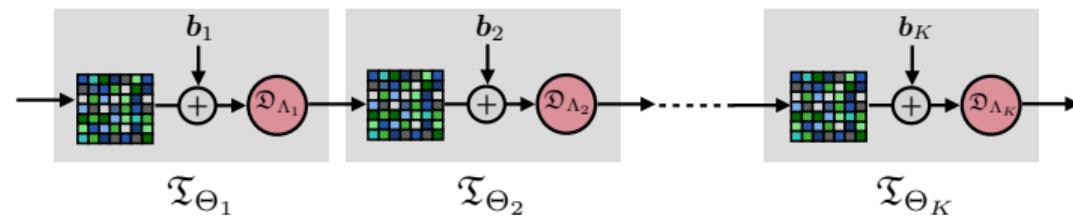


Nonlinearity



Dense matrix

**MLP**



# NN example: DnCNN

**DEFINITION:** DnCNN architecture takes the form

$$\mathcal{L}_{\Theta}^K = \mathfrak{T}_{\Theta_1} \circ \dots \circ \mathfrak{T}_{\Theta_K} \quad \text{with} \quad \mathfrak{T}_{\Theta_k} = \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + b_k)$$

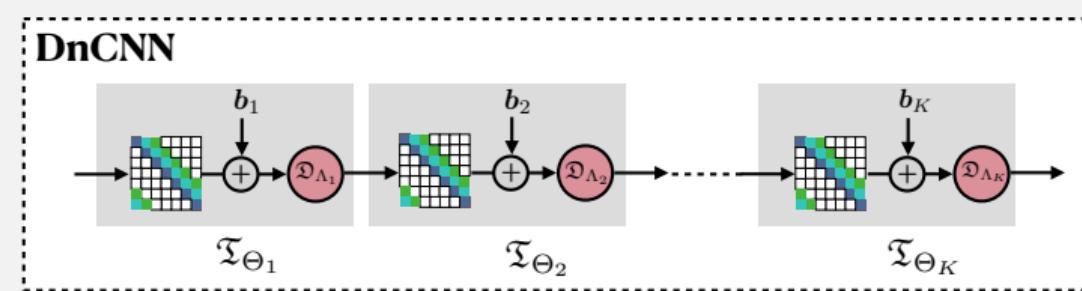
- $\mathfrak{D}_{\Lambda_1}$ : ReLU activation function
- $(\mathfrak{D}_{\Lambda_k})_{2 \leq k \leq K-1}$ : compositions of batch normalization and ReLU
- $\mathfrak{D}_{\Lambda_K}$ : identity
- $(\mathbf{L}_k)_{1 \leq k \leq K}$ : convolution operators

**REMARK:** Introduced as a denoising network [Zhang et al., 2017]

 Nonlinearity

 Dense matrix

 Convolution matrix



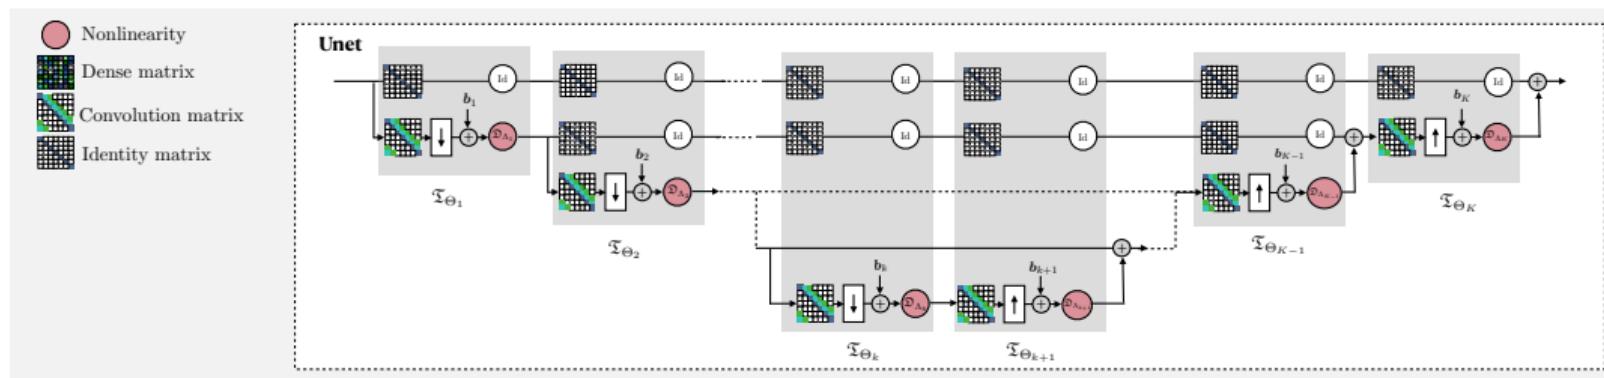
# NN example: U-Net

**DEFINITION:** Unet is an encoder-decoder architecture defined by

$$\mathcal{L}_{\Theta}^K = \mathfrak{T}_{\Theta_1} \circ \dots \circ \mathfrak{T}_{\Theta_K} \quad \text{with} \quad \mathfrak{T}_{\Theta_k} = \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + \mathbf{b}_k)$$

- $(\mathbf{L}_k)_{1 \leq k \leq K}$  are convolution operators
- $(\mathfrak{D}_{\Lambda_k})_{1 \leq k \leq K}$  are a combination of rectified linear units (ReLU) and max-pool/up-conv for upsampling/downsampling operations

**REMARK:** Developed for image segmentation [Ronneberger et al., 2015]



## NN example: U-Net

**DEFINITION:** Unet is an encoder-decoder architecture defined by

$$\mathcal{L}_{\Theta}^K = \mathfrak{T}_{\Theta_1} \circ \dots \circ \mathfrak{T}_{\Theta_K} \quad \text{with} \quad \mathfrak{T}_{\Theta_k} = \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + \mathbf{b}_k)$$

- $(\mathbf{L}_k)_{1 \leq k \leq K}$  are convolution operators
- $(\mathfrak{D}_{\Lambda_k})_{1 \leq k \leq K}$  are a combination of rectified linear units (ReLU) and max-pool/up-conv for upsampling/downsampling operations

**REMARK:** Unet architectures often use **skip-connections**:

- Considering a layer  $k$  where the variable needs to be “skipped”:

$$\mathfrak{T}_{\Theta_k}(\mathbf{x}) = \begin{pmatrix} \mathbf{x} \\ \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \mathbf{x} + \mathbf{b}_k) \end{pmatrix}$$

- Skip-connections and inertia/momentum share the same idea of reusing information from a past layer/iteration.

# Outline

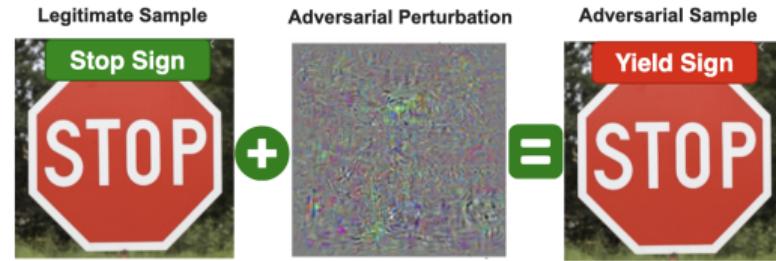
ROBUSTNESS: Lipschitz constant

LINK BETWEEN NEURAL NETWORKS AND PROXIMITY OPERATORS

- Activation functions and proximity operators
- Activation, potential and penalization operators

## THEORETICAL MOTIVATION: ROBUSTNESS OF NNs

# Robustness & Lipschitz constant

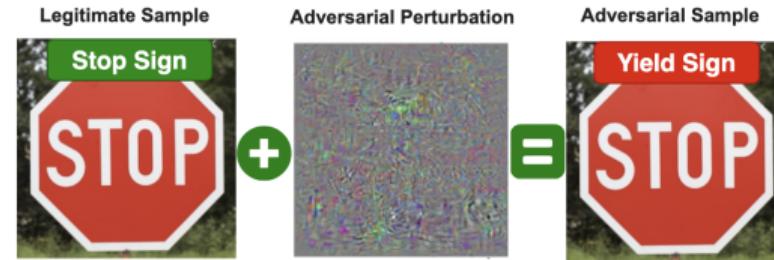


[ Ahmad et al. (2022)]

**QUESTION:** How to control robustness of NNs?

**IDEA:** Use Lipschitz constant [Hoffman et al. (2019)] [Jakubovitz and Giryes (2018)] [Pesquet et al. (2021)]

# Robustness & Lipschitz constant



[ Ahmad et al. (2022)]

**QUESTION:** How to control robustness of NNs?

**IDEA:** Use Lipschitz constant [Hoffman et al. (2019)] [Jakubovitz and Giryes (2018)] [Pesquet et al. (2021)]

**DEFINITION:** Let an input  $x \in \mathbb{R}^{N_0}$  and a perturbation  $\epsilon \in \mathbb{R}^{N_0}$ .

The error on the output can be bounded as:  $\|\mathcal{L}_\Theta^K(x + \epsilon) - \mathcal{L}_\Theta^K(x)\| \leq \omega \|\epsilon\|$

**REMARKS:** •  $\omega > 0$  is (by definition) a Lipschitz constant of  $\mathcal{L}_\Theta^K$

- $\omega$ : certificate of the network robustness

- Important and complementary to the performance in terms of estimation quality

# Robustness: Control of the Lipschitz constant

**DEFINITION:** Let an input  $\mathbf{x} \in \mathbb{R}^{N_0}$  and a perturbation  $\epsilon \in \mathbb{R}^{N_0}$ .

The error on the output can be bounded as:  $\|\mathfrak{L}_\Theta^K(\mathbf{x} + \epsilon) - \mathfrak{L}_\Theta^K(\mathbf{x})\| \leq \omega \|\epsilon\|$

- **LAYER-BY-LAYER CERTIFICATE:**

Let  $\mathfrak{L}_\Theta^K = \mathfrak{T}_{\Theta_1} \circ \dots \circ \mathfrak{T}_{\Theta_K}$ , with  $\mathfrak{T}_{\Theta_k} = \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + \mathbf{b}_k)$ . If  $\mathfrak{D}_{\Lambda_k}$  is nonexpansive, then a

Lipschitz constant for  $\mathfrak{L}_\Theta^K$  is  $\omega = \prod_{k=1}^K \|\mathbf{L}_k\|_S$

- **GLOBAL CERTIFICATE:**

If  $\mathfrak{L}_\Theta^K$  is differentiable, its Lipschitz constant is given by  $\omega = \sup_{\mathbf{x} \in \mathbb{R}^{N_0}} \|\nabla \mathfrak{L}_\Theta^K(\mathbf{x})\|_S$  where

$$\nabla \mathfrak{L}_\Theta^K(\mathbf{x}) = \left( \frac{\partial_i \mathfrak{L}_\Theta^K(\mathbf{x})}{\partial \mathbf{x}} \right)_{1 \leq i, j \leq N_0} \text{ is the Jacobian of } \mathfrak{L}_\Theta^K \text{ at } \mathbf{x} \in \mathbb{R}^{N_0}$$

# Layer-by-layer certificate

- **LAYER-BY-LAYER CERTIFICATE:**

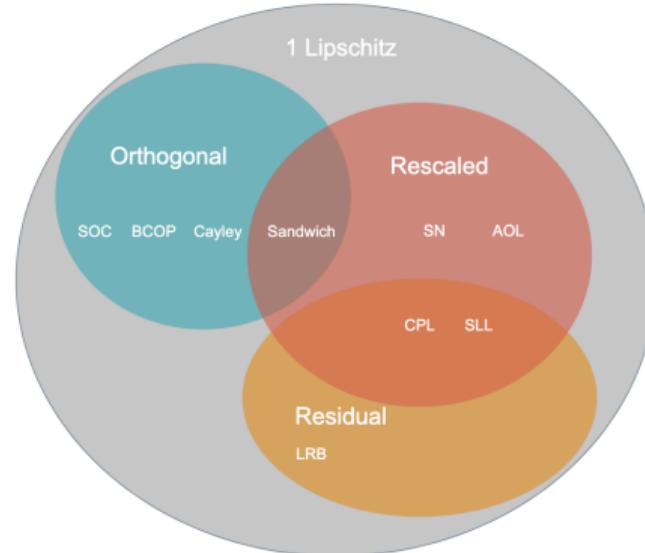
Let  $\mathfrak{L}_\Theta^K = \mathfrak{T}_{\Theta_1} \circ \dots \circ \mathfrak{T}_{\Theta_K}$ , with  $\mathfrak{T}_{\Theta_k} = \mathfrak{D}_{\Lambda_k}(\mathbf{L}_k \cdot + \mathbf{b}_k)$ . If  $\mathfrak{D}_{\Lambda_k}$  is nonexpansive, then a

Lipschitz constant for  $\mathfrak{L}_\Theta^K$  is  $\omega = \prod_{k=1}^K \|\mathbf{L}_k\|_S$

## REMARKS:

- Bound not tight
- Bounds can often be overly pessimistic in practice, thus providing limited insights into the robustness of the network [Combettes, Pesquet (2020)][Neacsu et al. (2024)].

# Some methods to control the Lipschitz constant of NN layers



[Delattre (2025)]

**Fig. 2.1.:** Venn diagram illustrating the categorization of Lipschitz layers, including scaled layers, orthogonal layers, and residual layers.

# Global certificate

- **GLOBAL CERTIFICATE:**

If  $\mathfrak{L}_\Theta^K$  is differentiable, its Lipschitz constant is given by  $\omega = \sup_{\mathbf{x} \in \mathbb{R}^{N_0}} \|\nabla \mathfrak{L}_\Theta^K(\mathbf{x})\|_S$

## REMARKS:

- Impossible to compute over all images in  $\mathbb{R}^N \rightsquigarrow$  Known to be an NP-hard problem
- Restrict the computation to a set of images  $(\mathbf{x}_i)_{i \in \mathbb{I}}$  [Pesquet et al. 2021], i.e.

$$\omega \approx \max_{(\mathbf{x}_i)_{i \in \mathbb{I}}} \|\nabla \mathfrak{L}_\Theta^K(\mathbf{x}_i)\|_S$$

- $\rightsquigarrow$  Underestimation of the Lipschitz constant.
- $\rightsquigarrow$  Computed using power iterations coupled with automatic differentiation.
- $\rightsquigarrow$  Used for constraining the value of  $\omega$  [Pesquet et al. 2021] (*see section PnP*)
- $\rightsquigarrow$  Used to compare robustness of different architectures [Le et al. 2023] (*see section unfolded*)

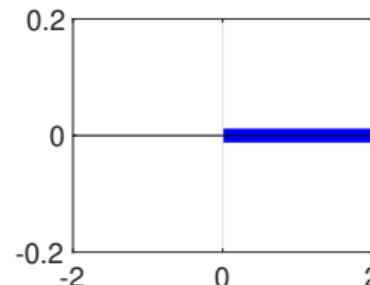
## LINKS BETWEEN NNs AND PROXIMITY OPERATORS

# Activation functions and proximity operators

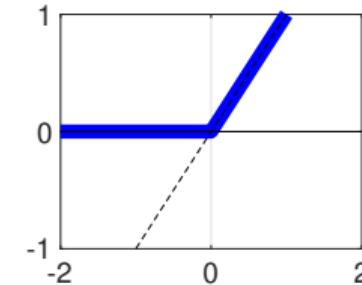
**REMARK:** Most activation functions used in imaging are proximity operators  
[Combettes & Pesquet, 2020]

**EXAMPLES:** ReLu

Potential  
 $\varphi(x) = \iota_{[0, +\infty)}$



Activation/ proximity operator  
 $\mathfrak{D}_\Lambda(x) = \text{prox}_\varphi(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$



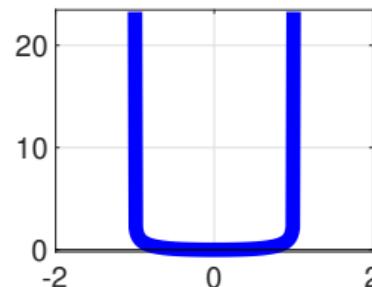
# Activation functions and proximity operators

**REMARK:** Most activation functions used in imaging are proximity operators  
[Combettes & Pesquet, 2020]

**EXAMPLES:** Arctangent

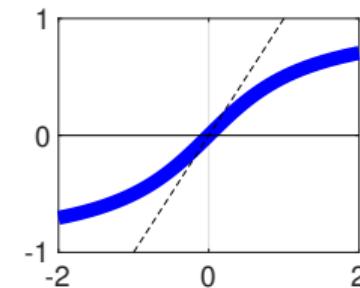
Potential

$$\varphi(x) = \begin{cases} -\frac{2}{\pi} \ln \left( \cos \left( \frac{\pi x}{2} \right) \right) - \frac{x^2}{2} & \text{if } |x| < 1 \\ +\infty & \text{otherwise} \end{cases}$$



Activation/ proximity operator

$$\mathfrak{D}_\Lambda(x) = \text{prox}_\varphi(x) = \frac{2 \arctan(x)}{\pi}$$



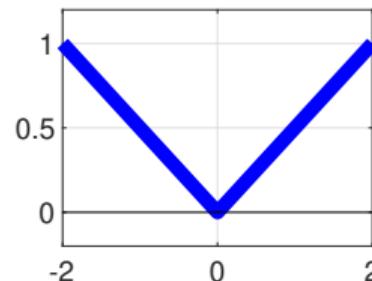
# Activation functions and proximity operators

**REMARK:** Most activation functions used in imaging are proximity operators  
[Combettes & Pesquet, 2020]

**EXAMPLES:** Soft thresholding

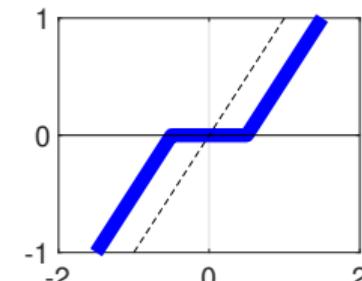
Potential

$$\varphi(x) = \lambda|x|$$



Activation/ proximity operator

$$\mathfrak{D}_\Lambda(x) = \text{prox}_\varphi(x) = \text{sign}(x) \max\{|x| - \lambda, 0\}$$



# Link between activation functions and proximity operators

**PROPOSITION** [Combettes & Pesquet, 2020] [Gharbi et al., 2023]

Let  $\varrho: \mathcal{H} \rightarrow \mathcal{H}$  be a function deriving from a potential, i.e., there exists a **differentiable** function  $\psi: \mathcal{H} \rightarrow \mathbb{R}$  such that  $\varrho = \nabla \psi$ .

- Let  $\beta \in (0, +\infty)$  and let  $\mu \in (-1, +\infty)$  be such that  $\beta = 1/(1 + \mu)$ .  
Then  $\varrho$  is the proximity operator of a  $\mu$ -convex function  $\varphi$  (i.e.,  $\varrho = \text{prox}_\varphi$ )  
 $\Leftrightarrow \psi$  is convex with a  $\beta$ - Lipschitz continuous gradient

**REMINDER:** •  $\varphi$  is  $\mu$ -convex if  $\varphi - \mu \frac{\|\cdot\|^2}{2}$  is convex

# Link between activation functions and proximity operators

**PROPOSITION** [Combettes & Pesquet, 2020] [Gharbi et al., 2023]

Let  $\varrho: \mathcal{H} \rightarrow \mathcal{H}$  be a function deriving from a potential, i.e., there exists a **differentiable** function  $\psi: \mathcal{H} \rightarrow \mathbb{R}$  such that  $\varrho = \nabla \psi$ .

- Let  $\beta \in (0, +\infty)$  and let  $\mu \in (-1, +\infty)$  be such that  $\beta = 1/(1 + \mu)$ .

Then  $\varrho$  is the proximity operator of a  $\mu$ -convex function  $\varphi$  (i.e.,  $\varrho = \text{prox}_\varphi$ )

$\Leftrightarrow \psi$  is convex with a  $\beta$ - Lipschitz continuous gradient

- $\varrho$  is  $\alpha$ -averaged with  $\alpha \in [1/2, 1]$

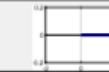
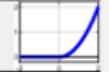
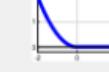
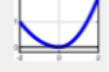
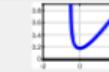
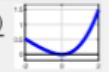
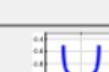
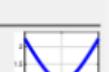
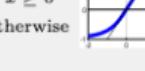
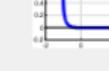
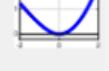
$\Leftrightarrow$  it is a  $(2\alpha)$  over-relaxation of the proximity operator of a convex function  $\varphi$ , that is

$$(\forall x \in \mathcal{H}) \quad \varrho(x) = x + 2\alpha(\text{prox}_\varphi x - x)$$

**REMINDER:** •  $\varphi$  is  $\mu$ -convex if  $\varphi - \mu \frac{\|\cdot\|^2}{2}$  is convex

•  $\varrho$  is  $\alpha$ -averaged if  $(1 - \alpha^{-1})\text{Id} + \alpha^{-1}\varrho$  is nonexpansive

# Activation functions, potentials and penalizations: Some examples

| Name                            | Prox/Activation<br>$x \in \mathbb{R} \mapsto \varrho(x) = \text{prox}_\varphi(x) = \psi'(x)$     | Potential<br>$x \in \mathbb{R} \mapsto \varphi(x)$                                | Penalization<br>$x \in \mathbb{R} \mapsto \psi(x)$  |   |  |   |
|---------------------------------|--|---|---|---|--|---|
| ReLU                            | $\begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$                         |  | $\iota_{[0, +\infty)}$  |  | $\begin{cases} \frac{x^2}{2} & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$                     |  |
| Parametric ReLU                 | $\begin{cases} x & \text{if } x > 0 \\ \lambda x & \text{otherwise} \end{cases}$                 |  | $\begin{cases} 0 & \text{if } x > 0 \\ \frac{(1/\lambda - 1)x^2}{2} & \text{otherwise} \end{cases}$   |  | $\begin{cases} \frac{x^2}{2} & \text{if } x > 0 \\ \lambda \frac{x^2}{2} & \text{otherwise} \end{cases}$ |  |
| Bent identity                   | $\frac{x + \sqrt{x^2 + 1} - 1}{2}$   |  | $\begin{cases} \frac{\pi}{2} - \frac{\ln(x + 1/2)}{4} & \text{if } x > -1/2 \\ +\infty & \text{otherwise} \end{cases}$  |  | $\frac{x^2 + x\sqrt{x^2 + 1} - 2x + \text{arcsinh}(x)}{4}$   |  |
| Inverse square root             | $\frac{x}{\sqrt{1+x^2}}$   |  | $\begin{cases} -\frac{x^2}{2} - \sqrt{1-x^2} & \text{if }  x  \leq 1 \\ +\infty & \text{otherwise} \end{cases}$   |  | $\sqrt{1+x^2}$   |  |
| Inverse square root linear unit | $\begin{cases} x & \text{if } x \geq 0 \\ \frac{x}{\sqrt{1+x^2}} & \text{otherwise} \end{cases}$ |  | $\begin{cases} 0 & \text{if } x \geq 0 \\ 1 - \frac{x^2}{2} - \sqrt{1-x^2} & \text{if } -1 \leq x < 0 \\ +\infty & \text{otherwise} \end{cases}$                                  |  | $\begin{cases} \frac{x^2}{2} & \text{if } x \geq 0 \\ \sqrt{1+x^2} - 1 & \text{otherwise} \end{cases}$   |  |
| Arctangent                      | $\frac{2 \arctan(x)}{\pi}$   |  | $\begin{cases} -\frac{2}{\pi} \ln \left( \cos \left( \frac{\pi x}{2} \right) \right) - \frac{x^2}{2} & \text{if }  x  < 1 \\ +\infty & \text{otherwise} \end{cases}$              |  | $\frac{2 \arctan(x)x - \ln( x^2 + 1 )}{\pi}$   |  |
| Unimodal sigmoid                | $\frac{1}{1 + \exp(-x)} - \frac{1}{2}$   |  | $\begin{cases} (\frac{1}{2} + x) \ln(\frac{1}{2} + x) & \text{if }  x  < 1/2 \\ -\frac{1}{2}(x^2 + \frac{1}{4}) & \text{if }  x  = 1/2 \\ +\infty & \text{otherwise} \end{cases}$ |  | $\frac{\pi}{2} + \ln(1 + \exp(-x))$  |  |

# Activation functions & proximity operators: Weakly convex case

**PROPOSITION** [Hurault et al., 2022]

Let  $\beta \in [0, 1]$  and let  $\psi: \mathcal{H} \rightarrow \mathbb{R}$  be a differentiable function with a  $\beta$ -Lipschitz gradient. Then, there exists a  $\beta/(\beta + 1)$ -weakly convex function  $\varphi: \mathcal{H} \rightarrow (-\infty, +\infty]$ , such that

$$\text{Id} - \nabla\psi = \text{prox}_\varphi$$

**REMARK:** If  $\varphi$  is convex, proper and l.s.c., then  $\psi$  is the Moreau envelope of  $\varphi$  defined as

$$(\forall \tilde{\mathbf{x}} \in \mathcal{H}) \quad \psi(\tilde{\mathbf{x}}) = \inf_{\mathbf{x} \in \mathcal{H}} \varphi(\mathbf{x}) + \frac{1}{2} \|\mathbf{x} - \tilde{\mathbf{x}}\|^2.$$

# Conclusion

